

```
In [28]: import pandas as pd

df = pd.read_excel('data/dataset.xlsx')
print(df.head())
```

	Date	Student	Study Hours	Sleep (hrs)	Exercise (mins)	\
0	01/06/2024	Asha	4	6	30	
1	01/06/2024	Ravi	3	7	20	
2	01/06/2024	Tina	1	8	10	
3	01/06/2024	Mike	2	6	20	
4	01/06/2024	Sara	1	6	25	

	Social Score	Screen Time (hrs)	Mood	Stress	Academic Score	Notes
0	5	3	6	6	83	Low energy
1	2	8	4	7	74	Felt anxious
2	5	3	10	2	71	NaN
3	2	6	8	2	77	NaN
4	4	7	8	3	76	Good focus

```
In [29]: print(df.info())           # Data types, missing values, column names
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  150 non-null   object
1   Student               150 non-null   object
2   Study Hours           150 non-null   int64
3   Sleep (hrs)           150 non-null   int64
4   Exercise (mins)       150 non-null   int64
5   Social Score           150 non-null   int64
6   Screen Time (hrs)     150 non-null   int64
7   Mood                  150 non-null   int64
8   Stress                150 non-null   int64
9   Academic Score        150 non-null   int64
10  Notes                 120 non-null   object
dtypes: int64(8), object(3)
memory usage: 13.0+ KB
None
```

```
In [30]: print(df.describe())
```

	Study Hours	Sleep (hrs)	Exercise (mins)	Social Score \
count	150.000000	150.000000	150.000000	150.000000
mean	2.846667	6.993333	17.000000	3.133333
std	1.379313	1.426013	9.796589	1.445505
min	1.000000	5.000000	0.000000	1.000000
25%	2.000000	6.000000	10.000000	2.000000
50%	3.000000	7.000000	20.000000	3.000000
75%	4.000000	8.000000	25.000000	4.000000
max	5.000000	9.000000	30.000000	5.000000

	Screen Time (hrs)	Mood	Stress	Academic Score
count	150.000000	150.000000	150.000000	150.000000
mean	5.580000	7.000000	4.140000	76.833333
std	1.810743	2.158175	1.996742	6.115682
min	3.000000	2.000000	2.000000	64.000000
25%	4.000000	6.000000	2.000000	72.000000
50%	6.000000	7.000000	4.000000	76.000000
75%	7.000000	9.000000	5.000000	81.000000
max	8.000000	10.000000	9.000000	92.000000

```
In [31]: print(df.isnull().sum())
```

```
Date          0
Student        0
Study Hours    0
Sleep (hrs)    0
Exercise (mins) 0
Social Score   0
Screen Time (hrs) 0
Mood           0
Stress         0
Academic Score 0
Notes          30
dtype: int64
```

```
In [32]: #Lets explore asha's record
asha_df = df[df['Student'] == 'Asha']
print(asha_df.head())

# Days with high stress
high_stress_df = df[df['Stress'] >= 7]
```

	Date	Student	Study Hours	Sleep (hrs)	Exercise (mins)	\
0	01/06/2024	Asha	4	6	30	
5	02/06/2024	Asha	4	5	15	
10	03/06/2024	Asha	2	8	10	
15	04/06/2024	Asha	2	7	10	
20	05/06/2024	Asha	3	8	30	

	Social Score	Screen Time (hrs)	Mood	Stress	Academic Score	\
0	5	3	6	6	83	
5	2	4	6	5	72	
10	5	5	7	4	71	
15	5	6	10	2	70	
20	2	7	8	3	81	

	Notes
0	Low energy
5	Good focus
10	Good focus
15	Felt anxious
20	NaN

```
In [33]: grouped = df.groupby('Student').agg({'Mood': 'mean', 'Stress': 'mean'})
print(grouped)
```

	Mood	Stress
Student		
Asha	6.966667	4.166667
Mike	6.766667	4.233333
Ravi	7.133333	4.100000
Sara	7.000000	4.066667
Tina	7.133333	4.133333

```
In [34]: df.select_dtypes(include='number').corr()
```

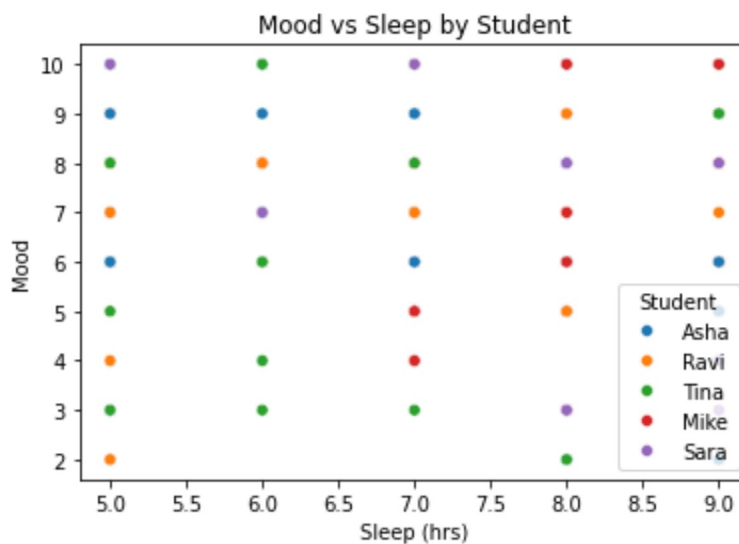
Out[34]:

	Study Hours	Sleep (hrs)	Exercise (mins)	Social Score	Screen Time (hrs)	Mood	Stress	Acad s
Study Hours	1.000000	0.006301	0.057615	-0.026705	0.086903	-0.056364	0.102884	0.71
Sleep (hrs)	0.006301	1.000000	0.176312	0.049273	-0.019286	0.030530	-0.004384	0.14
Exercise (mins)	0.057615	0.176312	1.000000	0.021327	0.178198	-0.238075	0.230904	0.13
Social Score	-0.026705	0.049273	0.021327	1.000000	0.011282	0.352818	-0.339023	0.00
Screen Time (hrs)	0.086903	-0.019286	0.178198	0.011282	1.000000	-0.166587	0.170440	-0.01
Mood	-0.056364	0.030530	-0.238075	0.352818	-0.166587	1.000000	-0.932893	0.04
Stress	0.102884	-0.004384	0.230904	-0.339023	0.170440	-0.932893	1.000000	0.01
Academic Score	0.712213	0.149937	0.132183	0.000253	-0.013636	0.041188	0.011267	1.00

In [35]:

```
import matplotlib.pyplot as plt
import seaborn as sns

sns.scatterplot(data=df, x='Sleep (hrs)', y='Mood', hue='Student')
plt.title('Mood vs Sleep by Student')
plt.show()
```



In [37]:

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
import numpy as np
import pandas as pd
```

```
def regression_analysis(df, features, target):  
    """  
    Fits a linear regression model and prints evaluation metrics and feature  
    importances.  
  
    Args:  
        df: DataFrame with your data  
        features: list of column names to use as predictors  
        target: string, the column to predict  
    """  
    X = df[features]  
    y = df[target]  
  
    # Handle missing values  
    X = X.fillna(X.mean())  
    y = y.fillna(y.mean())  
  
    # Train/test split  
    X_train, X_test, y_train, y_test = train_test_split(X, y,  
                                                         test_size=0.2,  
                                                         random_state=42)  
  
    # Fit model  
    model = LinearRegression()  
    model.fit(X_train, y_train)  
  
    # Predict  
    y_pred = model.predict(X_test)  
  
    # Evaluate  
    r2 = r2_score(y_test, y_pred)  
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))  
  
    # Feature importance (coefficients)  
    importance = pd.Series(model.coef_, index=features)  
  
    # Print results  
    print(f"Results for predicting {target}:")  
    print(f"R2 Score: {r2:.3f}")  
    print(f"RMSE: {rmse:.3f}")  
    print("\nFeature importances (effect on prediction):")  
    print(importance.sort_values(ascending=False))  
    print("-" * 40)  
  
    return model, importance, y_test, y_pred  
  
# Example usage for Mood:  
features_reg = ['Sleep (hrs)', 'Study Hours', 'Stress', 'Exercise (mins)', 'Screen  
Time (hrs)', 'Social Score']  
features_clf = ['Sleep (hrs)', 'Study Hours', 'Exercise (mins)', 'Screen Time  
(hrs)', 'Social Score', 'Mood']  
  
mood_model, importance, y_test, y_pred = regression_analysis(df, features_reg,  
                                                             'Mood')
```

```
# Example usage for Academic Score:
academic_model, importance, y_test, y_pred = regression_analysis(df, features_reg,
'Academic Score')
```

Results for predicting Mood:

R2 Score: 0.868

RMSE: 0.824

Feature importances (effect on prediction):

```
Sleep (hrs)          0.058492
Social Score         0.055681
Study Hours          0.050077
Screen Time (hrs)    0.015102
Exercise (mins)      -0.007887
Stress               -1.004148
dtype: float64
```

Results for predicting Academic Score:

R2 Score: 0.639

RMSE: 4.202

Feature importances (effect on prediction):

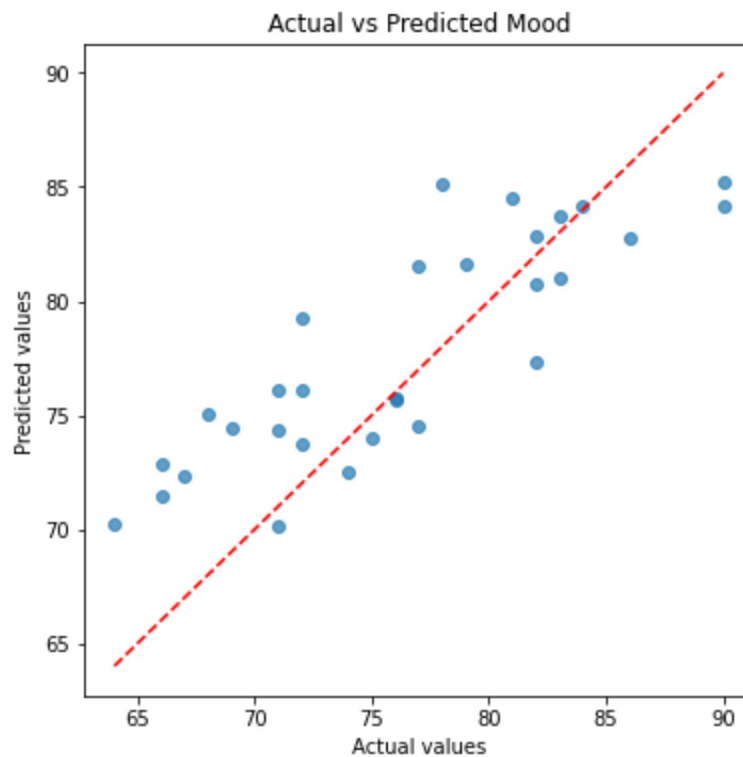
```
Study Hours          3.024389
Sleep (hrs)          0.595487
Exercise (mins)      0.052765
Social Score         -0.096717
Stress               -0.114289
Screen Time (hrs)    -0.316670
dtype: float64
```

In [38]: `import matplotlib.pyplot as plt`

```
def plot_predictions(y_test, y_pred, title="Actual vs Predicted"):
    plt.figure(figsize=(6,6))
    plt.scatter(y_test, y_pred, alpha=0.7)
    plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--') #
45-degree line
    plt.xlabel("Actual values")
    plt.ylabel("Predicted values")
    plt.title(title)
    plt.show()
```

In [15]: `# After y_pred = model.predict(X_test)`

```
plot_predictions(y_test, y_pred, title="Actual vs Predicted Mood")
```



```
In [39]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix

# Add a column: High Stress Day (1 if Stress >= 7, else 0)
df['HighStress'] = (df['Stress'] >= 7).astype(int)
X = df[features_clf]
y = df['HighStress']

# Train/test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Fit classifier
stress_model = LogisticRegression()
stress_model.fit(X_train, y_train)
y_pred = stress_model.predict(X_test)

# Evaluate
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

Accuracy: 0.9333333333333333

Confusion Matrix:

```
[[24  0]
 [ 2  4]]
```

	precision	recall	f1-score	support
0	0.92	1.00	0.96	24
1	1.00	0.67	0.80	6
accuracy			0.93	30
macro avg	0.96	0.83	0.88	30
weighted avg	0.94	0.93	0.93	30

```
In [40]: importance = pd.Series(stress_model.coef_[0], index=features_clf)
print("Feature importance (positive: increases high stress risk):")
print(importance.sort_values(ascending=False))
```

Feature importance (positive: increases high stress risk):

```
Study Hours      0.704498
Exercise (mins)  0.005058
Screen Time (hrs) -0.101397
Social Score     -0.185977
Sleep (hrs)      -0.215006
Mood             -2.318610
dtype: float64
```

```
In [41]: import os
os.makedirs('models', exist_ok=True)
```

```
In [42]: import pickle

with open('models/mood_regressor.pkl', 'wb') as f:
    pickle.dump(mood_model, f)

with open('models/academic_regressor.pkl', 'wb') as f:
    pickle.dump(academic_model, f)

with open('models/stress_classifier.pkl', 'wb') as f:
    pickle.dump(stress_model, f)
```

```
In [22]: # Example: predict Mood for new data
sample = pd.DataFrame([[7, 4, 6, 20, 5, 3]], columns=features_reg)
predicted_mood = mood_model.predict(sample)
print(predicted_mood)
```

[5.11639811]

```
In [ ]:
```